

METPO: A Pragmatic Ontology for Microbial Ecophysiological Traits

Mark Miller

Lawrence Berkeley National Laboratory

ICBO 2025 • Tuesday, November 11

The Data Integration Challenge

Microbial trait data is exploding, but it's semantically inconsistent:

BacDive (nested JSON)

```
{  
  "culture_temp": {  
    "temperature":  
      "mesophilic"  
  }  
}
```

BactoTraits (multi-header CSV)

Trait_class	Trait_name	Unit
Temperature	Class	
	mesophile	

Madin et al. (flat table)

taxon, temp_pref
12345, mesophilic

Problem: Cannot query "all psychrophiles" across datasets without semantic normalization.

Solution: We need an ontology.

The Gap in the Ontology Landscape

We evaluated existing ontologies for microbial trait coverage:

Ontology	Last Updated	Status	Domain
MPO	2014	Unmaintained	Microbial Phenotypes
MicrO	2018	Unmaintained	Microbiology
MCO	2019	Unmaintained	Microbial Conditions
OMP	2024	Active	Microbe Phenotypes
PATO	Active	Active	Quality Attributes

Key findings:

- No single ontology provides comprehensive coverage for BacDive, BactoTraits, and Madin datasets
- Many domain-specific ontologies lack active maintenance
- Integration of multiple ontologies introduces significant complexity

Introducing METPO

Microbial Ecophysiological Trait and Phenotype Ontology

Design philosophy: Application-driven development

- **Focused:** 255 classes covering BacDive, BactoTraits, and Madin datasets
- **Lightweight:** Purpose-built for KG-Microbe knowledge graph
- **Modern:** ODK, ROBOT, OAK tooling
- **Maintained:** Active development for DOE CMM project (\$850K FY26)

Core principle: Build ontologies that serve real applications, driven by actual data integration needs.

METPO's Modern & Pragmatic Design

1. Captures Real-World Variations

METPO synonyms handle messy data:

```
METPO:1000615 a owl:Class ;
  rdfs:label "mesophilic" ;
  oboInOwl:hasExactSynonym
    "mesophile" ;
  oboInOwl:hasExactSynonym
    "temperature: mesophilic" ;
  oboInOwl:hasExactSynonym
    "temp_class: mesophile" .
```

Enables semantic normalization across all three data sources.

2. LLM-Assisted Curation

- Semantic search pipeline:
 - OLS + BioPortal + Names4Life embeddings
 - ~3,000 METPO mappings to external ontologies
 - SSSOM format (skos:exactMatch, closeMatch, relatedMatch, broadMatch)
- Modern development practices:
 - ODK build system for consistency
 - ROBOT validation in CI/CD
 - Spreadsheet-based curation for accessibility

ICBO 2025 Theme: AI-enabled ontology development

METPO in Production: KG-Microbe Statistics

Real-world usage across three major datasets:

Dataset	Records	METPO Edges	Coverage	METPO Terms Used
BactoTraits	19,456 strains	87,688	100%	91
Madin et al.	172,324 taxa	45,851	39.7%	36
BacDive	196,169 nodes	120,266	7.3%	61
Combined	1.86M edges	253,805	13.6%	152 unique

Source: Direct analysis of KG-Microbe transformed data (TSV files)

All statistics traceable to primary sources in [/kg-microbe/data/transformed/](#)

METPO in Action: Powering KG-Microbe

Before: Heterogeneous Data

```
// BacDive  
{"culture_temp": "mesophilic"}  
  
// BactoTraits  
[blank][mesophile][blank]  
  
// Madin  
temperature_preference: mesophilic
```

Three formats, three field names, slight variations in values.

Cannot query across datasets.

After: Unified Knowledge Graph (KGX)

subject	predicate	object
NCBITaxon:100	biolink:has_phenotype	METPO:1000615
NCBITaxon:200	biolink:has_phenotype	METPO:1000615
NCBITaxon:300	biolink:has_phenotype	METPO:1000615

All variations normalized to single METPO CURIE.

Now queryable with SPARQL/Cypher:

```
SELECT ?taxon WHERE {  
    ?taxon biolink:has_phenotype METPO:1000615  
}
```

KG-Microbe: METPO as Semantic Backbone

Production knowledge graph for growth media prediction:

- **253,805 METPO edges** across 1.86M total edges
- **152 unique METPO terms** normalize traits from three heterogeneous datasets
- **Biolink model compliance** enables standard graph queries
- **Application:** CMM project - predicting organisms for REE biorecovery

Figure: METPO usage distribution across KG-Microbe datasets

Real Data Example: KG-Microbe Triples

From transformed data (BacDive → KGX):

```
# Nodes
METPO:1000615 biolink:PhenotypicQuality mesophilic
METPO:1000602 biolink:PhenotypicQuality aerobic
METPO:1000614 biolink:PhenotypicQuality psychrophilic
METPO:1000644 biolink:PhenotypicQuality heterotrophic

# Edges
NCBITaxon:100 biolink:has_phenotype METPO:1000615 R0:0002200 bacdive:17384
NCBITaxon:100 biolink:has_phenotype METPO:1000602 R0:0002200 bacdive:17385
NCBITaxon:200 biolink:has_phenotype METPO:1000614 R0:0002200 bacdive:23451
```

This is production data - METPO CURIEs in use, enabling ML predictions.

Future Work: Literature-Based Discovery

Expanding METPO through OntoGPT experiments:

Input: PubMed abstract

Methylovirgula ligni gen. nov., sp. nov., an obligately acidophilic, facultatively methylotrophic bacterium...

Two strains of Gram-negative, aerobic, non-pigmented, non-motile, rod-shaped bacteria were isolated...

Output: RDF triples (Turtle)

```
AUTO:Bw863
  has_phenotype METPO:1000143 . # ✓ Gram-negative
AUTO:Bw863
  has_phenotype METPO:1000008 . # ✓ aerobic
AUTO:Bw863
  has_phenotype AUTO:rod-shaped . # ✗ gap!
AUTO:Bw863
  has_phenotype METPO:1000181 . # ✓ mesophilic
```

Results across 10 PubMed abstracts:

- 32.5% phenotype grounding to METPO (25/77 terms)
- 100% chemical grounding to ChEBI (47/47 terms)
- 100% taxonomic grounding to NCBI Taxon (22/22 terms)

Data-Driven Ontology Development

Failed groundings drive iterative expansion:

1. Literature mining identifies phenotype: "rod-shaped"
2. OntoGPT cannot ground → produces AUTO:rod-shaped
3. Triggers curation workflow
4. Domain expert adds term to METPO
5. Future extractions ground successfully

This is METPO's living development model.

OntoGPT Grounding: Data-Driven Expansion

Structured data integration (Phase 1): Production-ready

- 253,805 METPO edges in KG-Microbe
- Three major datasets fully integrated
- METPO provides semantic normalization layer

Literature mining (Phase 2): Experimental validation

- 10 PubMed abstracts processed through OntoGPT
- Mixed grounding success identifies where METPO needs expansion
- 52 failed groundings → curation priorities for domain experts

Insight: Failed groundings are features, not bugs—they guide ontology development

METPO Hierarchy Example

Organized for our use case:

- Temperature preferences (psychrophilic → mesophilic → thermophilic)
- Oxygen requirements (aerobic, anaerobic, microaerophilic)
- Metabolic capabilities (heterotrophic, autotrophic, mixotrophic)
- Chemical interaction properties (65+ predicates for organism-chemical relationships)

Interoperability: METPO is a Good Citizen

SSSOM Mappings to external ontologies:

Two mapping strategies evaluated:

Strategy	Total Mappings	Match Types
Combined Relaxed	2,659 broadMatch 171 relatedMatch 113 closeMatch 65 exactMatch	Mixed precision
Optimized	2,590 relatedMatch 220 closeMatch 73 exactMatch	Higher precision

Integration approach: Pragmatic mappings enable interoperability while maintaining METPO's focused design

Availability:

- BioPortal: <https://bioportal.bioontology.org/ontologies/METPO>

Sustainability & Governance

How we ensure METPO won't be abandoned:

1. **Active funding:** DOE CMM project (\$850K FY26) - REE biorecovery research
2. **Real application:** KG-Microbe depends on METPO for production predictions
3. **Modern tooling:** ODK, ROBOT, OAK prevent bitrot
4. **Accessible contribution:** ROBOT-compatible spreadsheets, not complex OWL
5. **LLM-assisted workflows:** Scalable curation, undergraduate-accessible

Governance:

- PI: Marcin Joachimiak (LBNL)
- Driven by application needs (pragmatic, not theoretical)
- Community contributions via GitHub

ICBO 2025 Theme: Long-term sustainability

Conclusion

METPO demonstrates an application-driven approach to ontology development:

- ✓ **Production validation:** 253,805 edges across 1.86M-edge knowledge graph
- ✓ **Sustainability by design:** Active funding, real application dependencies, modern tooling
- ✓ **Transparent metrics:** All claims traceable to primary data sources
- ✓ **Pragmatic interoperability:** ~3,000 mappings to external ontologies via SSSOM
- ✓ **AI-assisted workflows:** LLM-powered curation and semantic search pipelines

Key insight: Purpose-built ontologies can effectively serve specialized domains when:

- Driven by concrete application needs (KG-Microbe)
- Supported by active maintenance (CMM project funding)
- Built with modern practices (ODK, ROBOT, spreadsheet curation)
- Validated against real data (BacDive, BactoTraits, Madin datasets)

Thank You

METPO Resources:

- GitHub: <https://github.com/berkeleybop/metpo>
- BioPortal: <https://bioportal.bioontology.org/ontologies/METPO>

KG-Microbe:

- GitHub: <https://github.com/Knowledge-Graph-Hub/kg-microbe>

Contact:

- Mark Miller, LBNL
- PI: Marcin Joachimiak

Acknowledgments:

DOE CMM Program • LBNL Biosciences Division • OntoGPT/OAK teams

Backup Slides

CMM Project Context

Advanced Biorecovery of Critical Minerals through AI/ML-Guided Design

Goal: Recover Rare Earth Elements (Nd, Pr, Dy) from e-waste using engineered microbes

KG-CMREE extends KG-Microbe:

- Adds REE bioaccumulation phenotypes
- Lanthanophore biosynthetic pathways
- Spectroscopic measurements
- High-throughput screening results

METPO's role: Annotates microbial traits critical for REE recovery

- pH tolerance, temperature tolerance, metal tolerance, metabolic capabilities
- Enables AI models to predict which microbes can thrive in REE recovery conditions

Technical Details: METPO Development

Source of truth: ROBOT-compatible spreadsheets (Google Sheets)

Build system: Ontology Development Kit (ODK)

Validation: ROBOT (0 errors, 318 minor whitespace warnings)

Semantic search: Custom pipeline

- OLS + BioPortal + Names4Life embeddings in SQLite
- openai_text-embedding-3-small model
- Generates SSSOM mappings

Statistics:

- 255 total terms
- 118 with definitions (46.3%)
- 158 mapped to 24 external ontologies (62%)
- ~250 classes in hierarchy

Design Rationale: Why Build METPO?

Coverage challenge:

- Existing ontologies lack comprehensive coverage for our three datasets
- Integration complexity increases non-linearly with number of source ontologies
- Unmaintained ontologies present long-term sustainability risks

Practical approach:

- Build focused ontology for immediate application needs
- Map to external ontologies for interoperability (~3,000 SSSOM mappings)
- Maintain lightweight hierarchy optimized for our use cases

Development efficiency:

- 255 terms covers core needs vs. integrating dozens of partial ontologies
- Spreadsheet-based curation accessible to domain experts
- ODK/ROBOT automation prevents technical debt

Comprehensive Ontology Survey: What We Evaluated

Systematic testing documented in PRIMARY SOURCES:

ChromaDB testing (verified via SQL queries):

- **39 ontologies embedded** with 778,496 total embeddings
- Top tested: CHEBI (221K), upheno (192K), GO (84K), OBA (73K), foodon (40K)
- **24 ontologies retained:** 20 OLS + 4 BioPortal-only
- **15 removed:** Including CHEBI (worst ROI despite size)

Final retained set (from ChromaDB SQL queries):

- **20 from OLS:** upheno, go, oba, flopo, micro, pato, envo, ecocore, eupath, phipo, mco, eco, omp, ohmi, cmpo, biolink, apo, pco, geo, exo
- **4 from BioPortal (verified via API):**
 - d3o (DSMZ Digital Diversity Ontology)
 - meo (Metagenome and Microbes Environmental Ontology)
 - miso (Microbial Isolation Source Ontology)

Why Not Import? Minimum Import Set Analysis

Q: Based on high-quality matches, how many ontologies would we need to import?

Analysis of SSSOM mappings (PRIMARY SOURCE: .sssom.tsv file):

Excellent matches (similarity ≥ 0.75 , distance < 0.25):

- 182 total excellent matches across 21 ontologies
- Top 8 ontologies provide 91% coverage:
 - i. micro (91 matches)
 - ii. upheno (19)
 - iii. mpo (16)
 - iv. n4l_merged (14)
 - v. oba (10)
 - vi. envo (4)
 - vii. biolink (4)
 - viii. flopo (3)

Semantic-SQL Ontology Registry

INCATools semantic-sql project maintains a comprehensive ontology registry:

Registry location:

```
/home/mark/gitrepos/semantic-sql/src/semsql/builder/registry/ontologies.yaml
```

Purpose: Provides standardized SQL views for 113 OBO and biomedical ontologies

- Includes URLs, build commands, prefix mappings for each ontology
- Used to generate SQLite databases for efficient querying
- Examples: GO, CHEBI, ENVO, PATO, MONDO, HP, and many more

METPO in semantic-sql:

- METPO available as `metpo.db` in semantic-sql directory
- Queryable via standard SQL interfaces
- Most ontologies we use are specialized/domain-specific